



**kubernetes**  
**in production**


Meetup DockerGrunn 10 April 2017

Herman Bos <[hbos@osso.nl](mailto:hbos@osso.nl)>  
Alex Boonstra <[aboonstra@osso.nl](mailto:aboonstra@osso.nl)>  
Jelle Prins <[jprins@osso.nl](mailto:jprins@osso.nl)>

# Plan

- Introducing OSSO
- Overview infrastructure setup
- Current k8s deployment
- Adding kubernetes to the mix
- Some important k8s concepts
- Demo cluster for tonight
- Protips
- Q&A

# Introducing OSSO



- specializes in open source {hosting,network,cloud} infrastructure
- {dev,netdev,sys,sec}ops

## Our customers:

- SaaS/cloud services, hosting providers and tech-startups

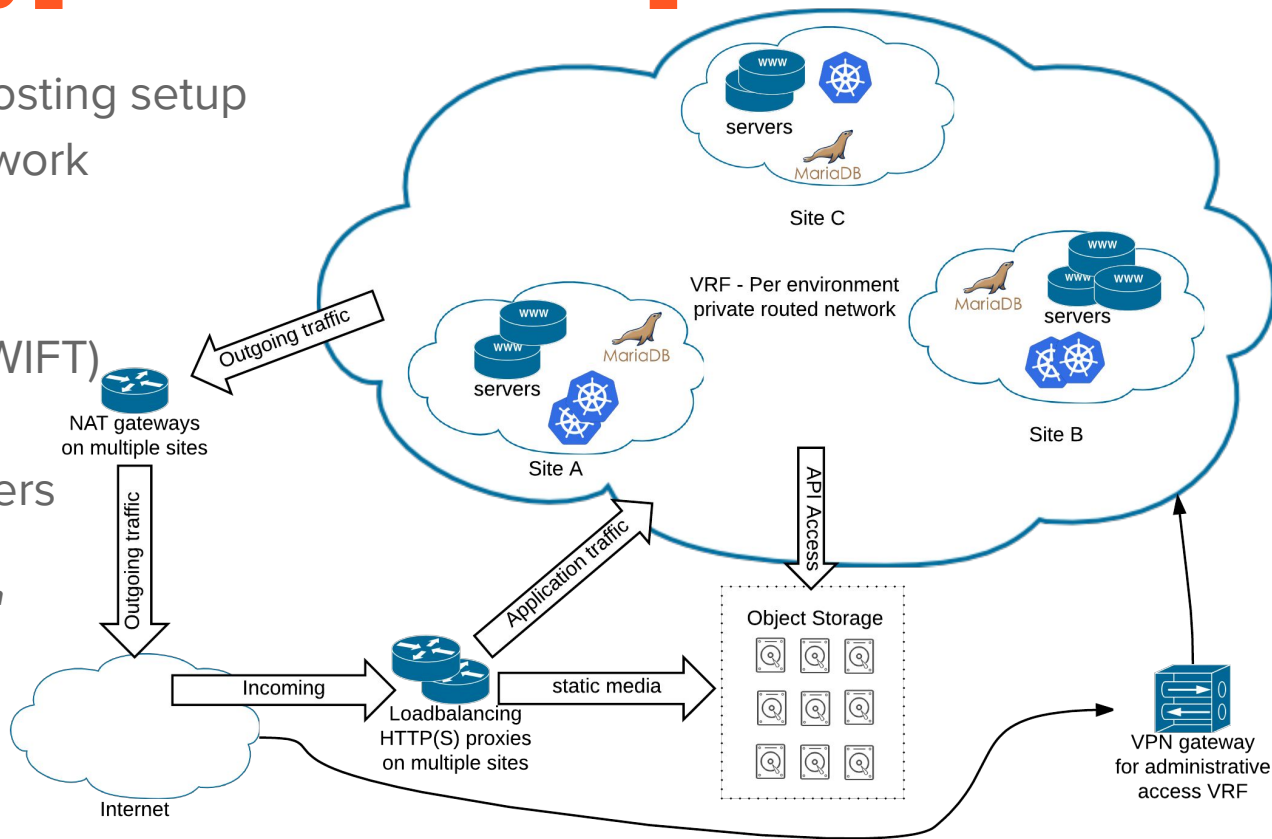
## We provide:

- 24/7 operations and support.
- Expertise, infrastructure development and so on.
- Multi datacenter high availability hosting infrastructure.

# Overview typical osso setup


- Multi datacenter hosting setup
- Private routed network
- Loadbalancer
- Nat gateway
- Object Storage (SWIFT)
- VPN
- {micro,virtual} servers

Mostly offered as "IaaS"



# Storage

- Object Storage for files (SWIFT, S3)
  - Cheap, no maintenance, high available fix.
- Other data in their respective services
  - MySQL, Redis, MongoDB, RabbitMQ, etc
  - Most common stuff can be deployed in a High availability configuration





# Our current deployment

- Current k8s deployments
  - 4 production clusters
  - 3-5 kubernetes nodes each currently
- We deploy k8s with saltstack on ubuntu 16.04lts



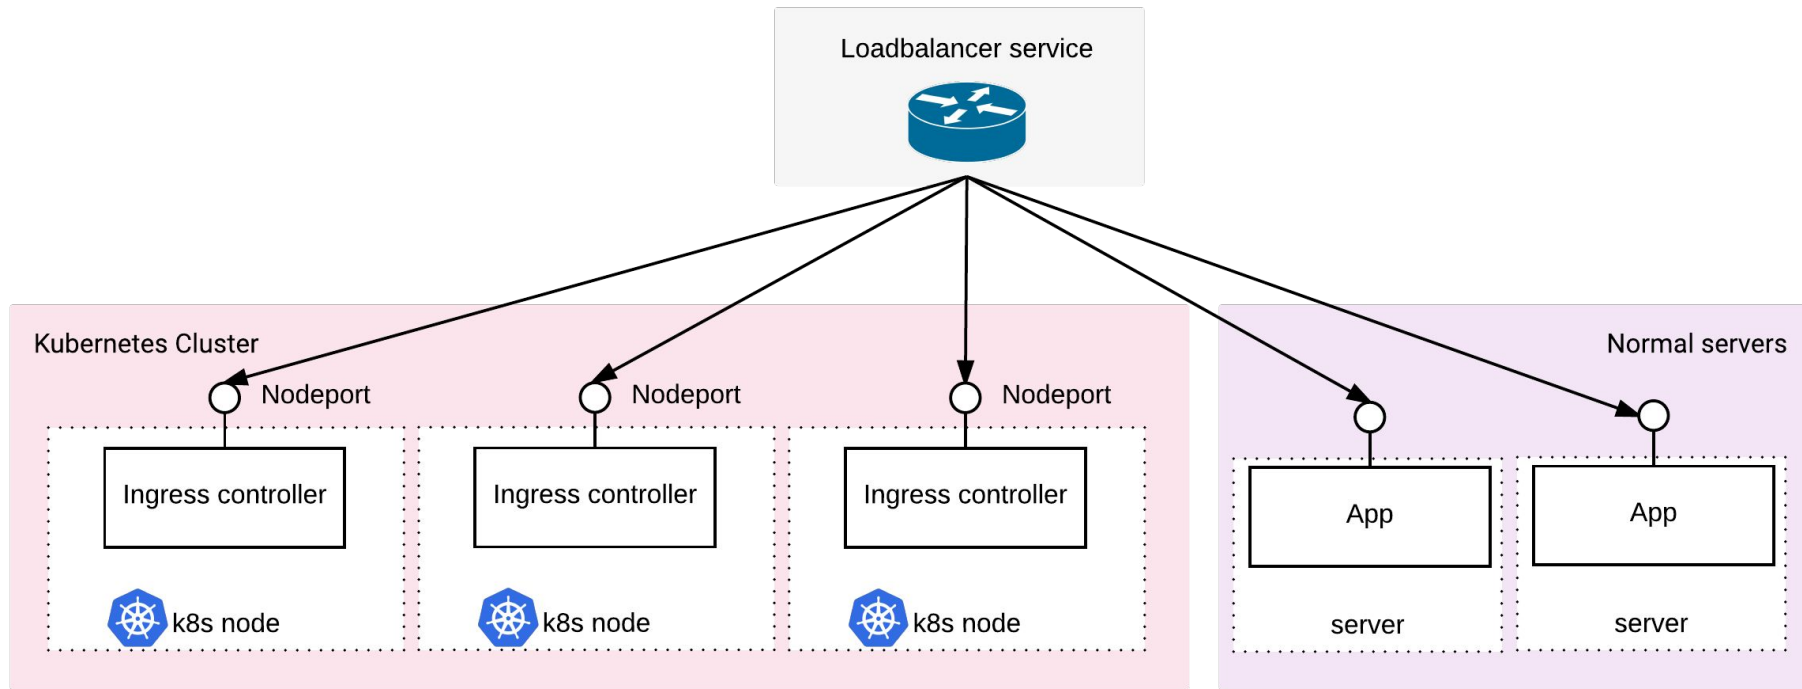
# Deploying k8s itself

- OSSO uses on Ubuntu 16.04LTS. Saltstack managed.
  - Same as we use in other environments (Ubuntu/Debian), a lot of available experience.
  - Can just apply all our existing configuration management, monitoring, etc.
  - Has public feed for security CVE status (which we use for gocollect)
- Initially started with Container Linux (CoreOS)
  - It is nice but we dropped it.
- Kubernetes distributions are emerging
  - e.g. CoreOS has Tectonic, Redhat has OpenShift, etc.

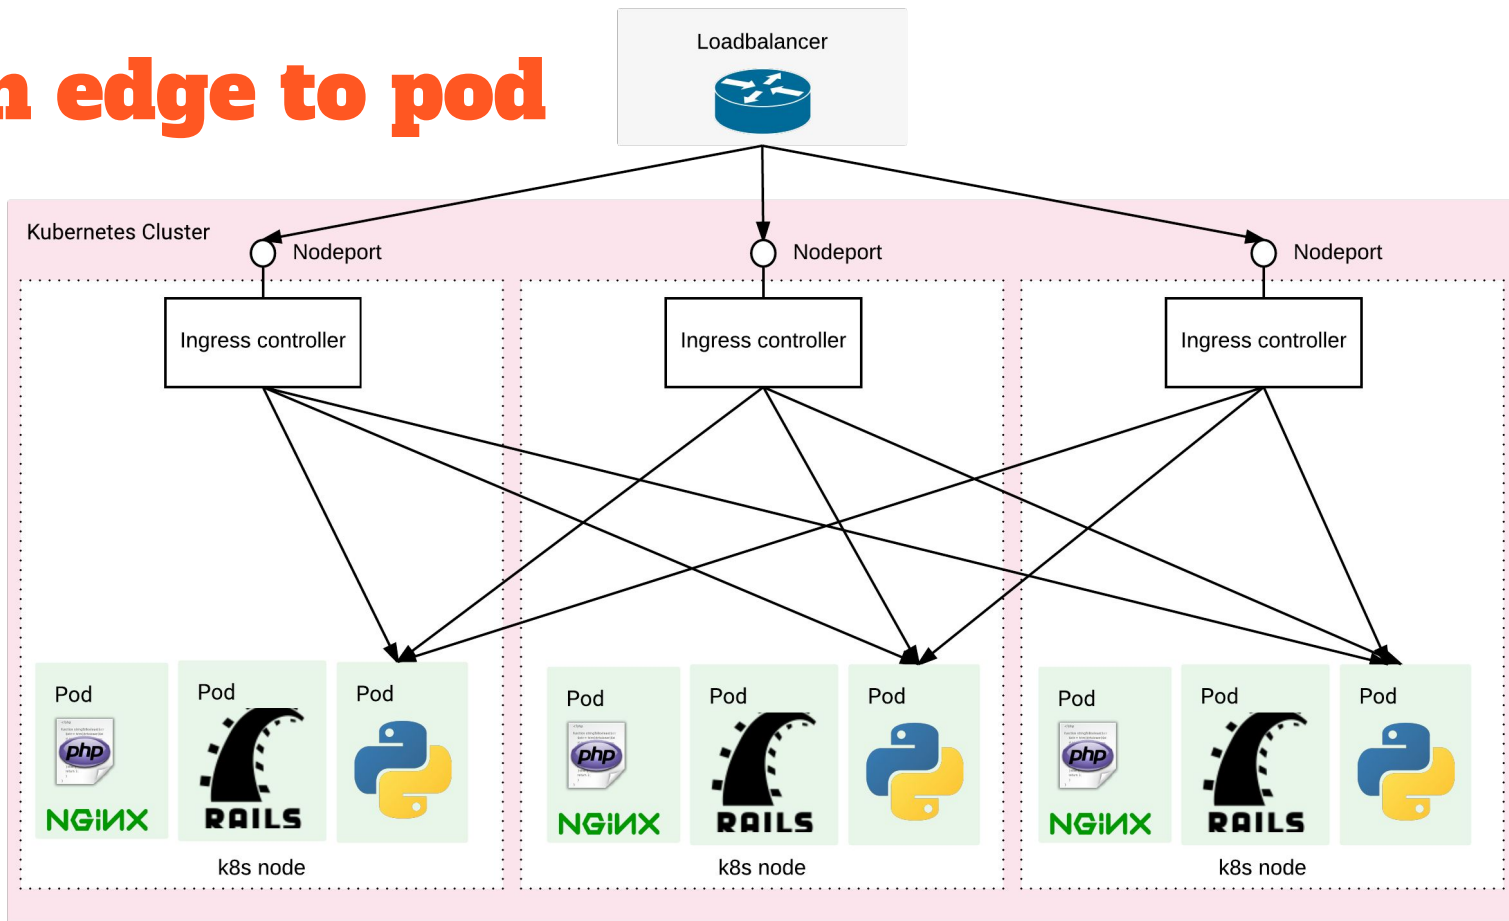


# Adding kubernetes to the mix

Run it next to your existing infrastructure and migrate your apps one by one (or even run them concurrently)



# From edge to pod



# K8s controllers

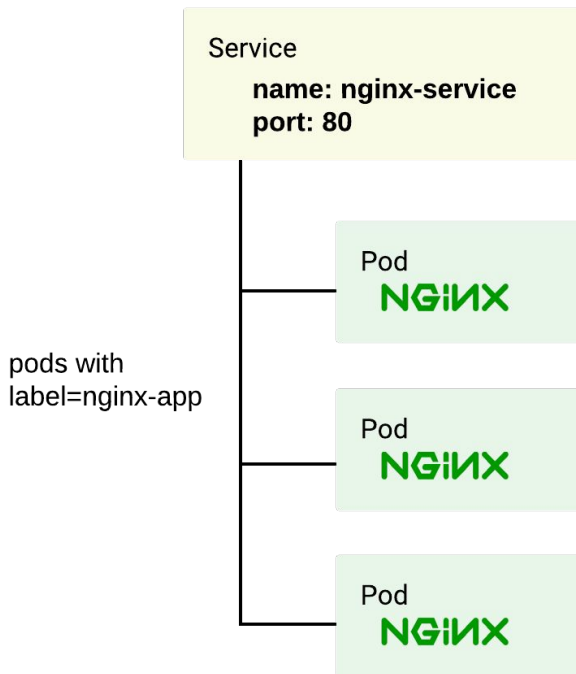
- **Replica Sets:** Controls how many pods are running
  - When to use: normally you use these indirectly through Deployments
- **Deployments:** Uses Replica Sets + deployment mechanics.
  - Use these for **stateless** applications (no persistent data etc).
- **StatefulSets:** For stateful services
  - Unique network identifiers, persistent storage, graceful deployment and scaling.
- **Daemon Sets:** For deploying a service on every node
  - For example ingress controllers, prometheus exporters, etc.

# Services, kube-dns, kube-proxy

- The correct way to connect to containers
- No need to keep track of container ip's
- Uses labels to target pods
- Map incoming port to a targetport inside a pod
- Works well together with kube-dns
- Service gets a virtual ip (which the dns will resolve to)
- Dns entry can be used from inside the application or from the ingress rules

# Services, kube-dns, kube-proxy

```
apiVersion: v1
kind: Service
metadata:
  namespace: awesome
  name: nginx-service
spec:
  selector:
    app: nginx-app
  ports:
    - port: 80
```



```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  namespace: awesome
  name: nginx-deployment
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: nginx-app
    spec:
      containers:
        - name: nginx-container
          image: ossobv/nginx:1.10.2
          ports:
            - containerPort: 80
```

# k8s ConfigMap and Secrets

- The right place to store your configuration parameters
- Secrets and ConfigMaps work similar

```
apiVersion: v1
data:
  tls.crt: #BASE64_DATA#
  tls.key: #BASE64_DATA#
kind: Secret
metadata:
  annotations:
    kubernetes.io/tls-acme: "true"
  creationTimestamp: 2017-04-10T13:02:11Z
  name: awesome-nginx-tls
  namespace: awesome
  resourceVersion: "605871"
  selfLink: /api/v1/namespaces/awesome/secrets/awesome-nginx-tls
  uid: e935817d-1ded-11e7-9fad-0cc47aeabeb1
type: kubernetes.io/tls
```


```
apiVersion: v1
data:
  default: |
    server {
      listen 80 default_server;
      root /var/www/html;
      index index.html;
      server_name _;
      location /healthz {
        return 200;
      }
      location / {
        try_files $uri $uri/ =404;
      }
    }
kind: ConfigMap
metadata:
  name: nginx-static-html-config
  namespace: awesome
```

# k8s.osso.ninja

Grab your ninja and deploy



Getting started instructions:  
<https://get.osso.ninja>





# Cluster k8s node specs

- E5-2630v4
- 128GB RAM
- 25Gbit/s network
- 240GB ssd raid1

32x 100G open networking switch

# k8s.osso.ninja demo cluster

Lb node

Node 3

Node 2

Node 1

# Pro tips

- Avoid shared file systems.
- Focus on getting your app deployments to K8s first.
  - Make then stateless if needed.
- Leave databases, legacy apps and other persistent data services for later.
  - And then use StatefulSets
- Run k8s cluster next to your current servers
  - Migrate apps/services one by one.
- Play around and get some experience
- Read the docs and ask questions

**Q&A**